

# Age and Gender Detection

*From Face Data*

Age: 1  
Gender: Male



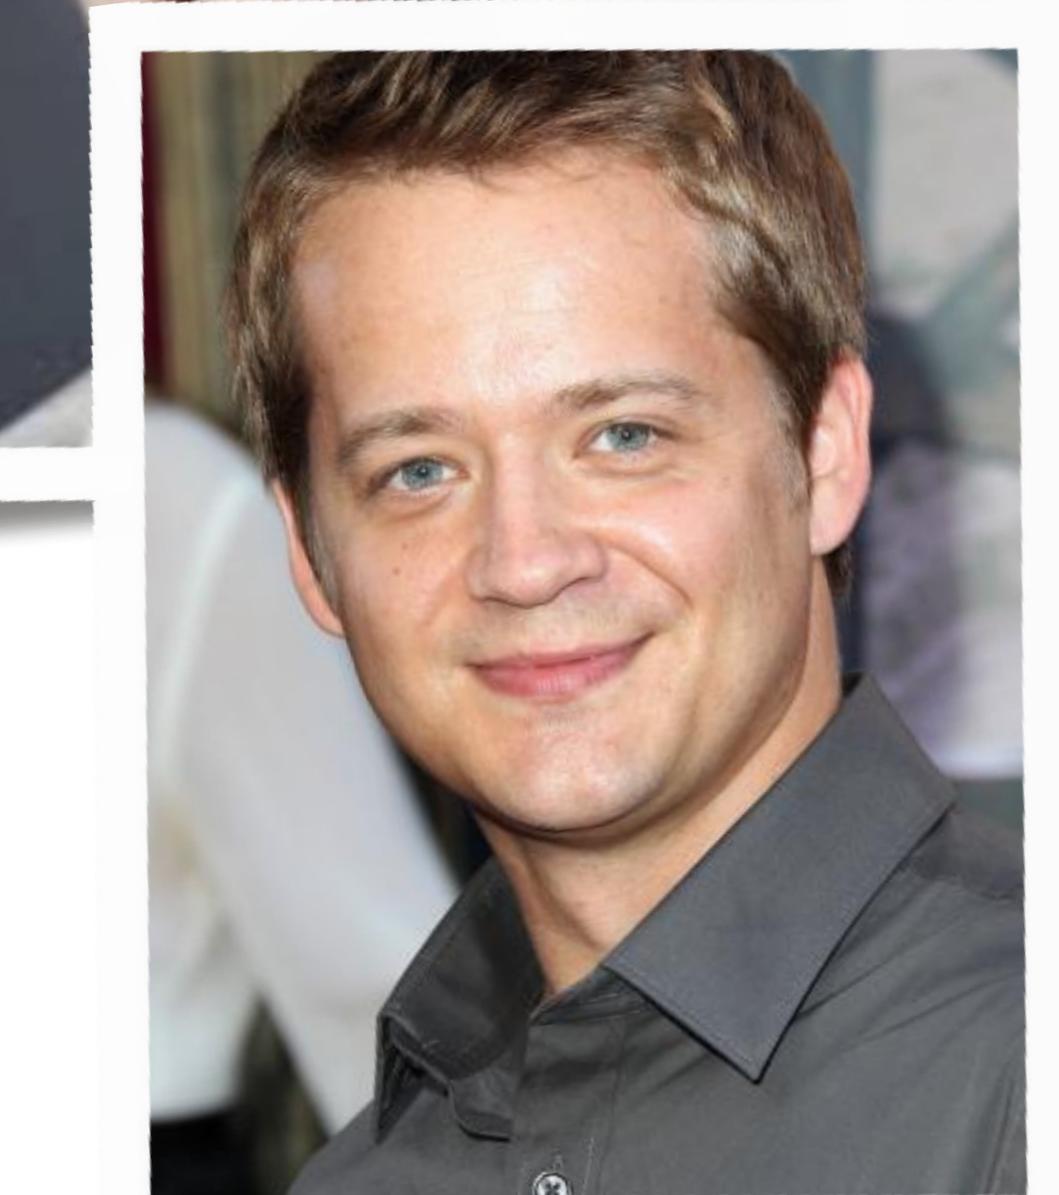
Hrayr Muradyan

# Why to identify gender and age?

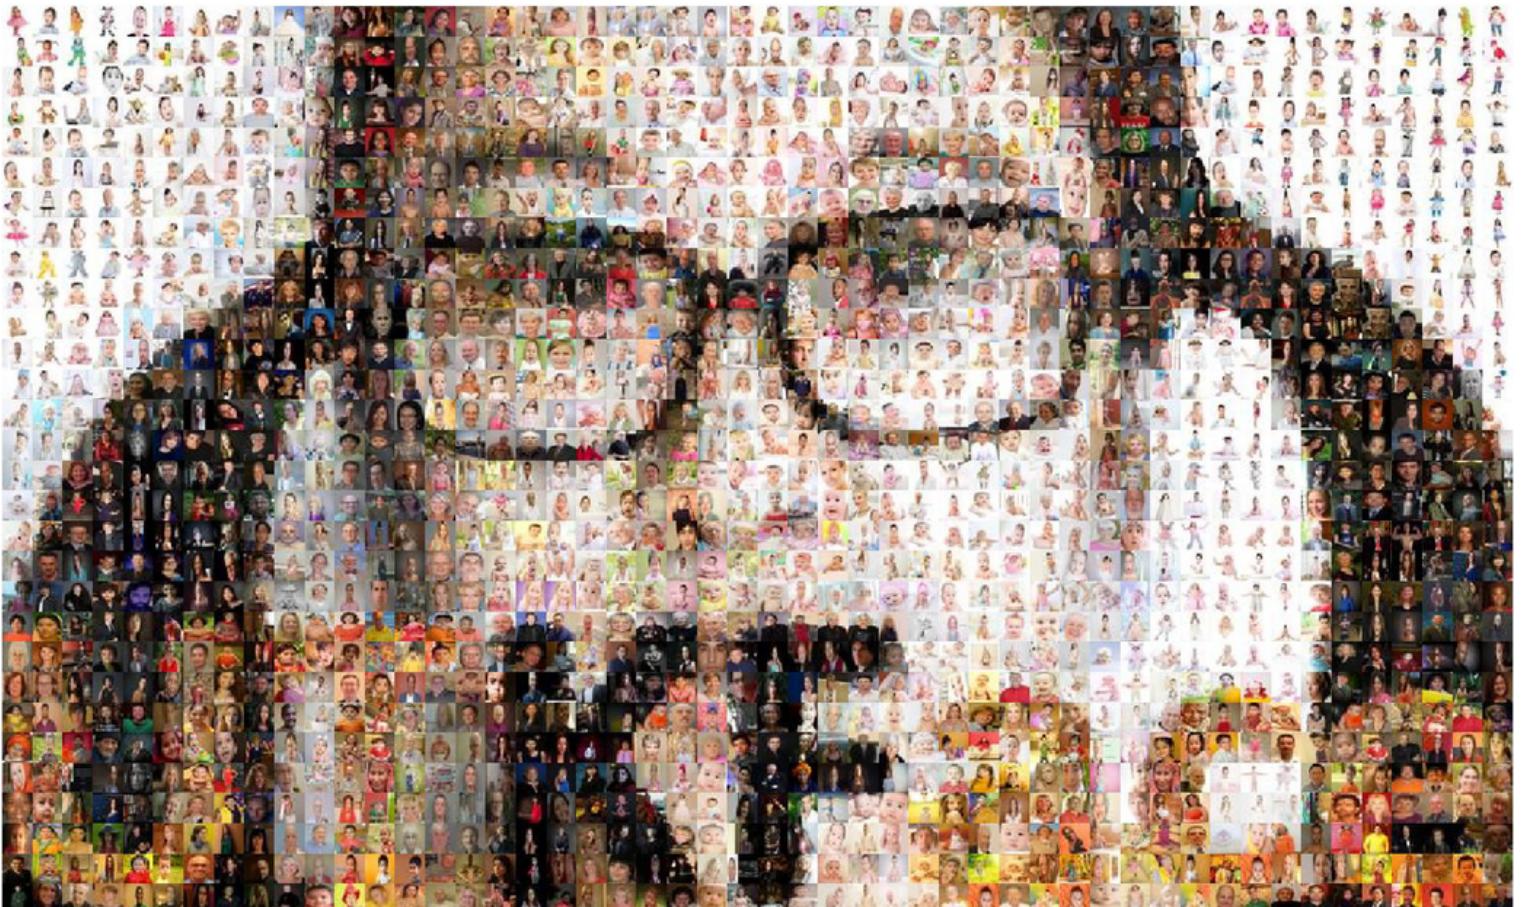
**Security:** In security systems, age and gender detection can aid in identifying individuals and understanding crowd compositions, which is crucial for safety and monitoring purposes.

**Healthcare and Elderly Care:** In healthcare, this technology can assist in monitoring elderly patients, tailoring care according to age and gender-specific needs.

**Age-Restricted Content Control:** It helps in ensuring that age-appropriate content is shown to users, especially crucial for online platforms hosting violate content.



# The Dataset



**UTKFace** dataset is a large-scale face dataset with long age span (range from 0 to 116 years old). The dataset consists of over **20,000** face images with annotations of age, gender, and ethnicity. The images cover large variation in pose, facial expression, illumination, occlusion, resolution, etc.

Source: <https://susanqq.github.io/UTKFace/>



# Annotations

**Each image information is contained in the filename:**

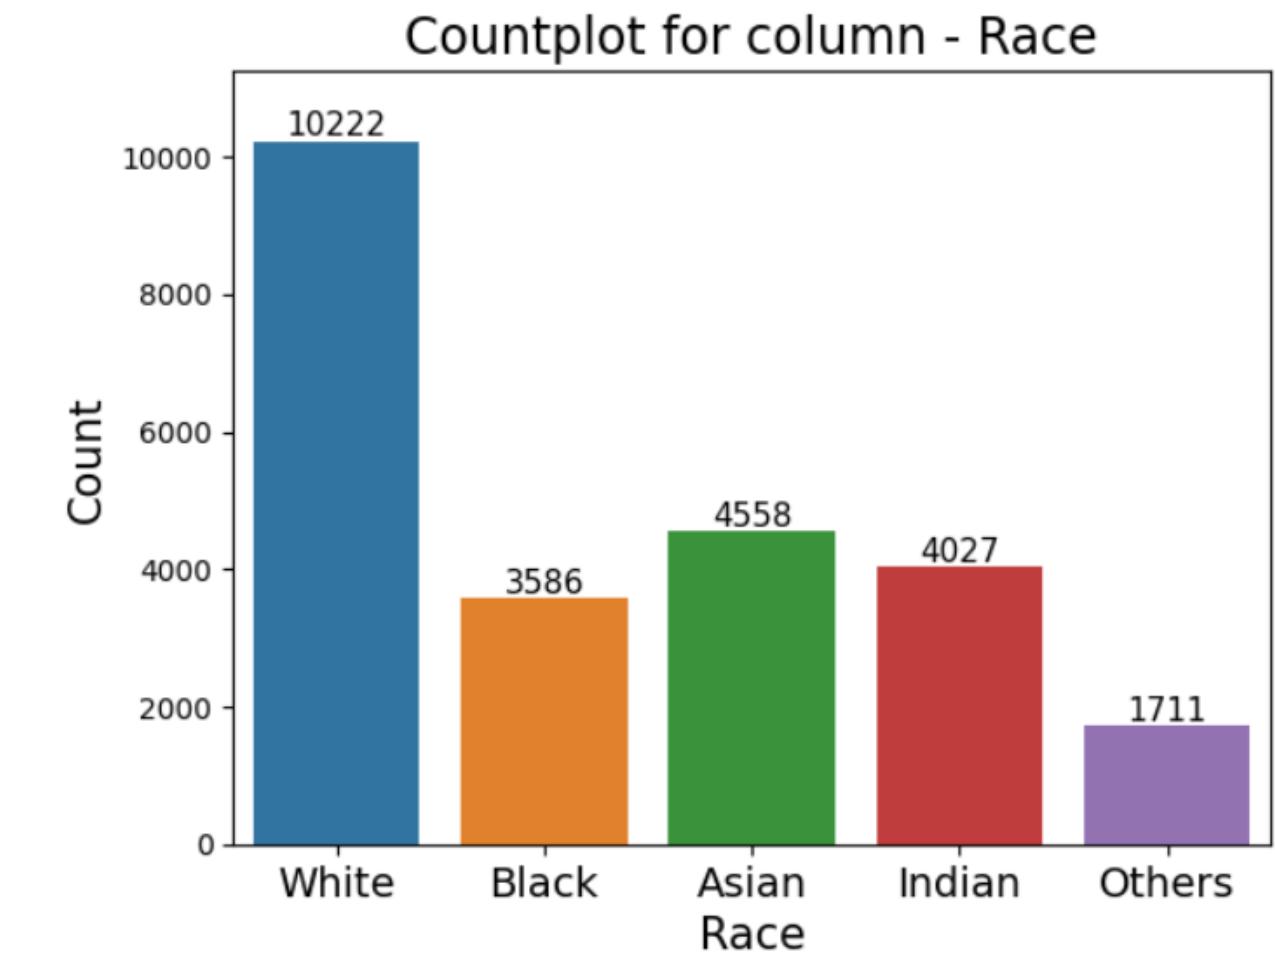
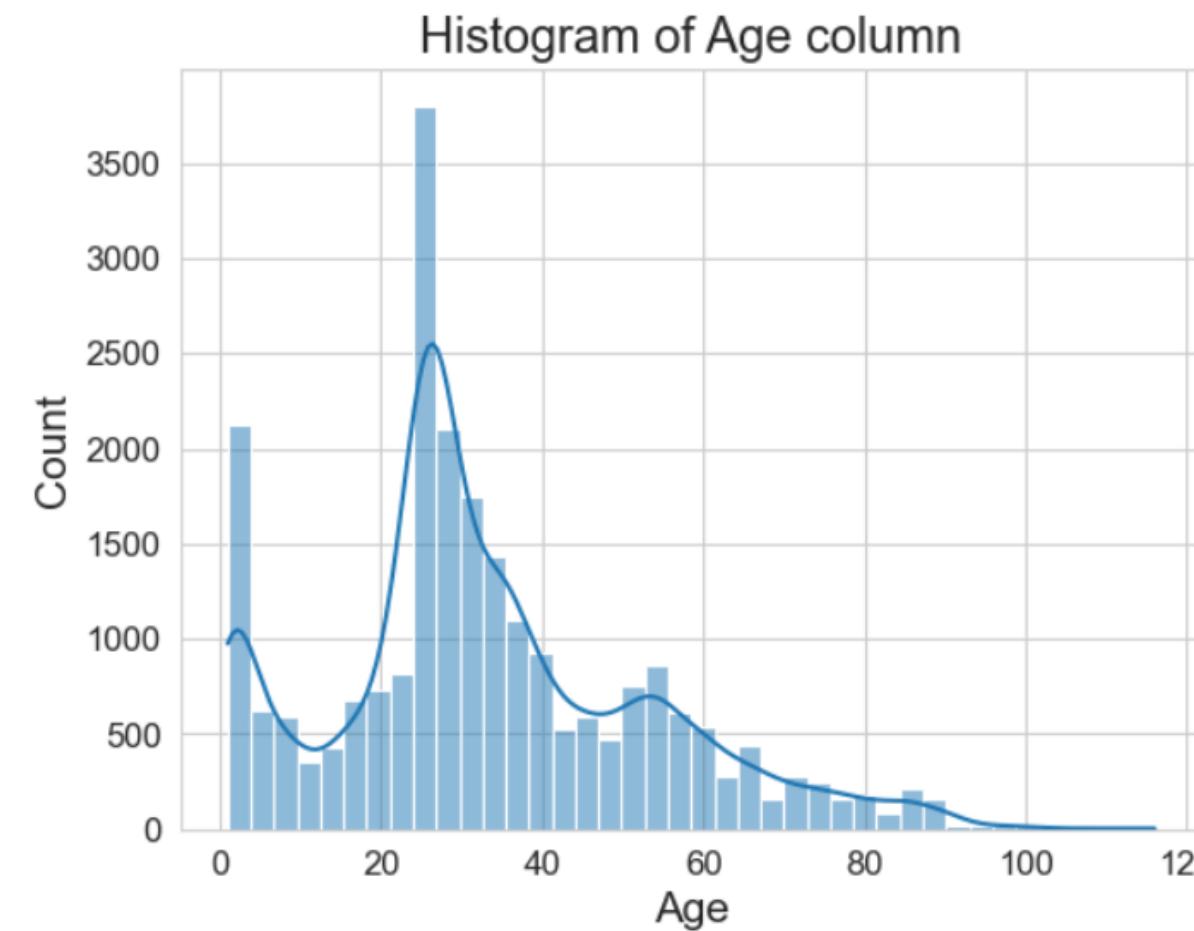
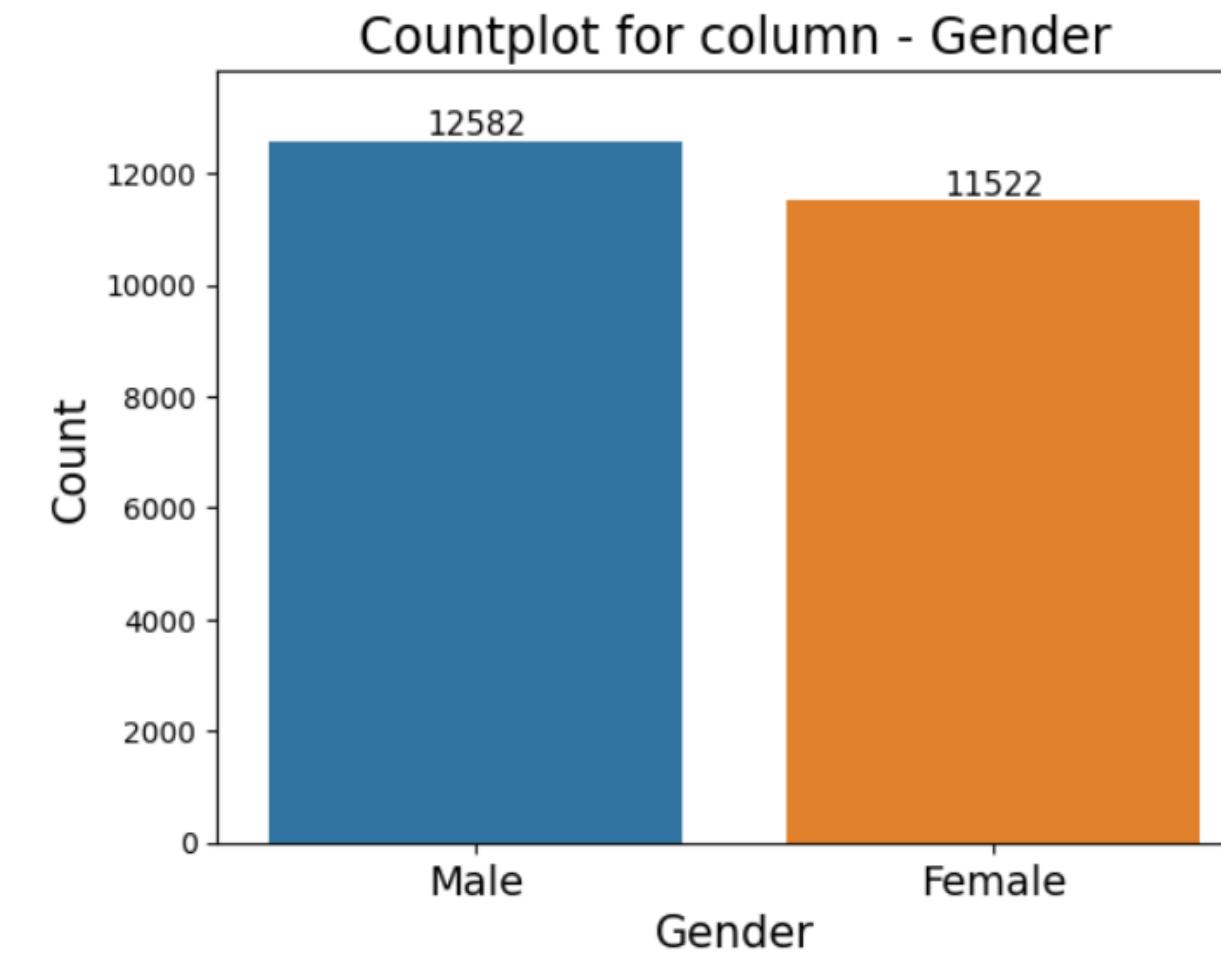
**[AGE]\_[GENDER]\_[RACE]\_[DATE&TIME].JPG**

- **[AGE]** is an integer from 0 to 116, indicating the age
- **[GENDER]** is either 0 (male) or 1 (female)
- **[RACE]** is an integer from 0 to 4, denoting White, Black, Asian, Indian, and Others



# Data Analysis

There are mostly middle-aged people, with approximately equal distribution of gender.



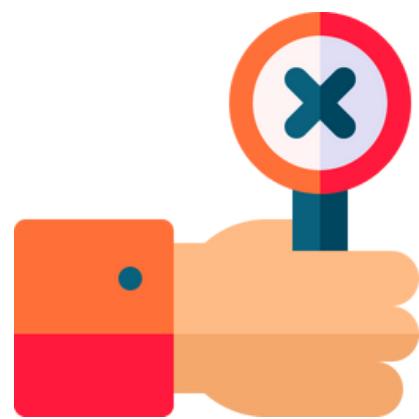
# Specifications

- Rotated images
- Bad quality images
- Two or more faces
- Side looking images
- Black and white images
- Wrong Images



⋮ ⋮  
⋮ ⋮  
⋮ ⋮  
⋮ ⋮

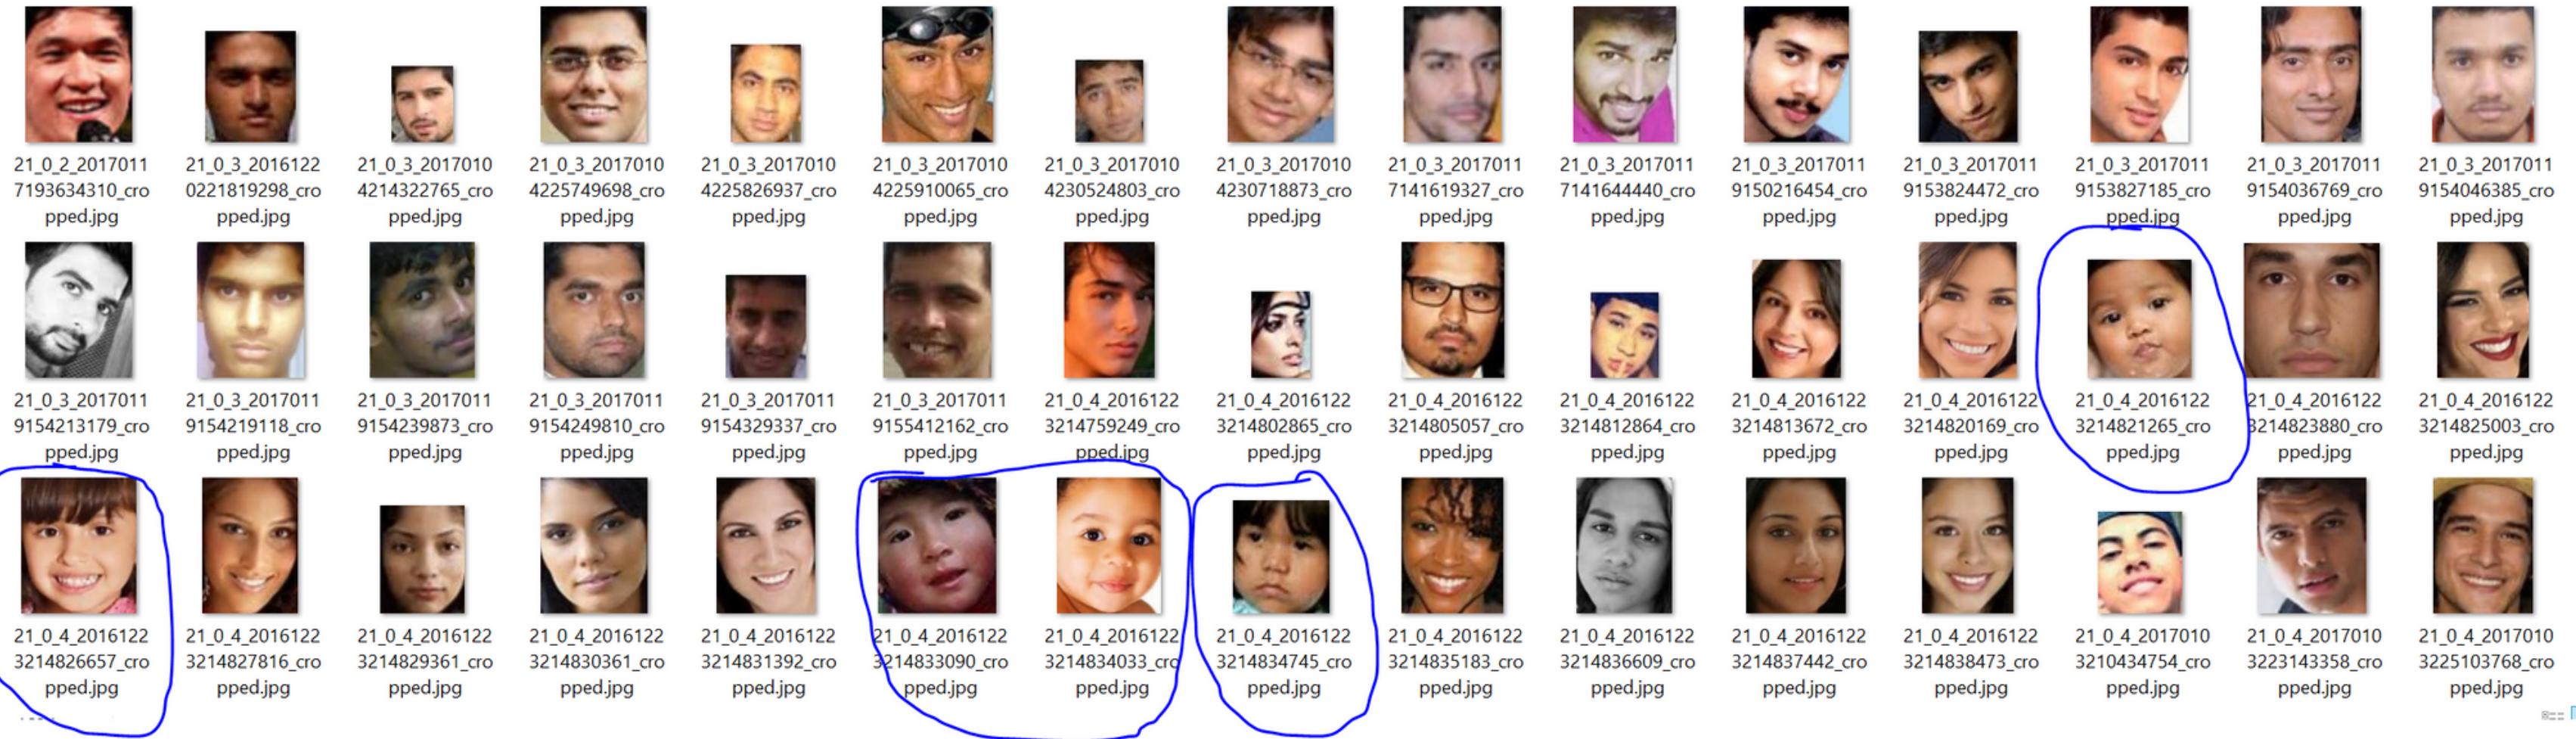
# Observations



# Wrong Data

Wrong data were deleted from the main dataset.

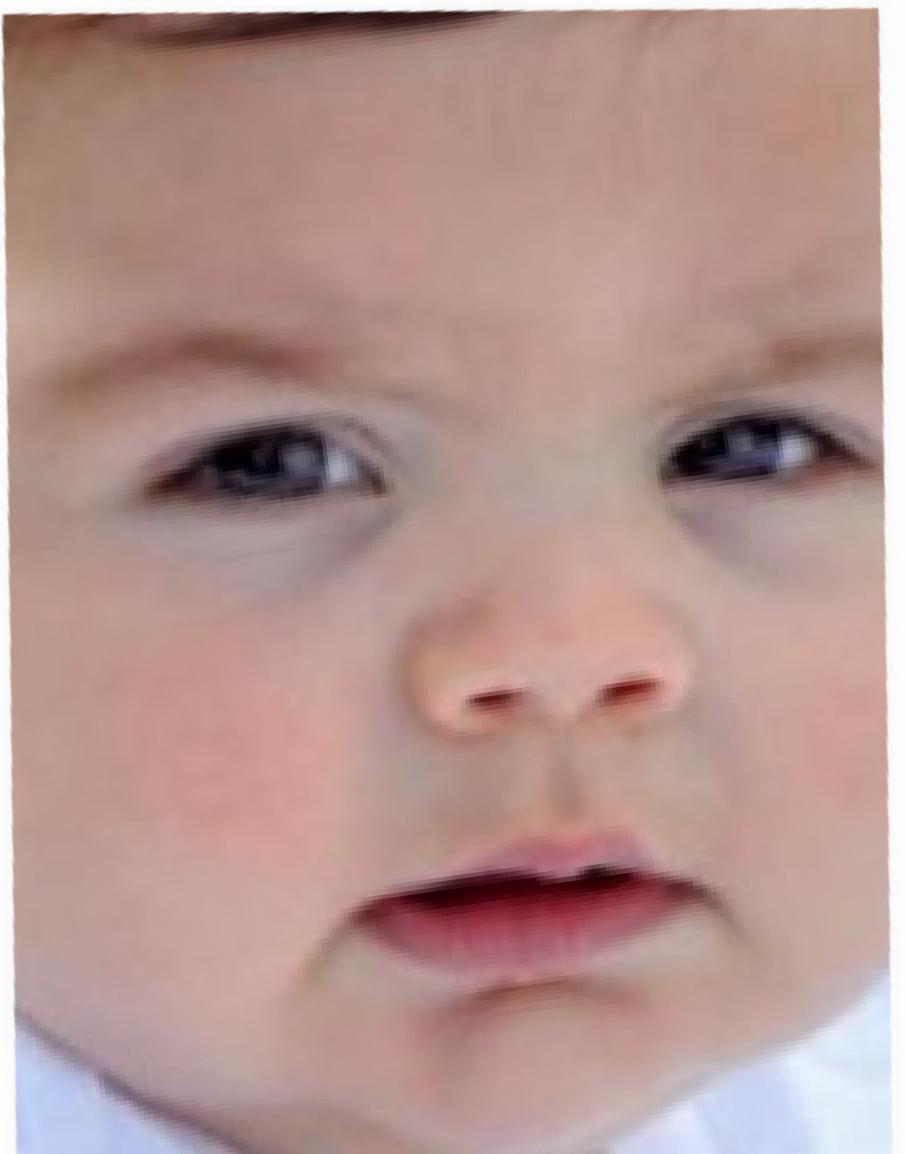
[images\39\\_1\\_20170116174525125.jpg](#)  
[images\61\\_1\\_20170109142408075.jpg](#)  
[images\61\\_3\\_20170109150557335.jpg](#)  
[images\61\\_3\\_20170109150557335.jpg](#)



# Crop Faces

The first step of the project is to **crop the faces** from the images, to concentrate the predictions more on the face of the person.

Two or more face images, or images where no face was found, **were all removed**.



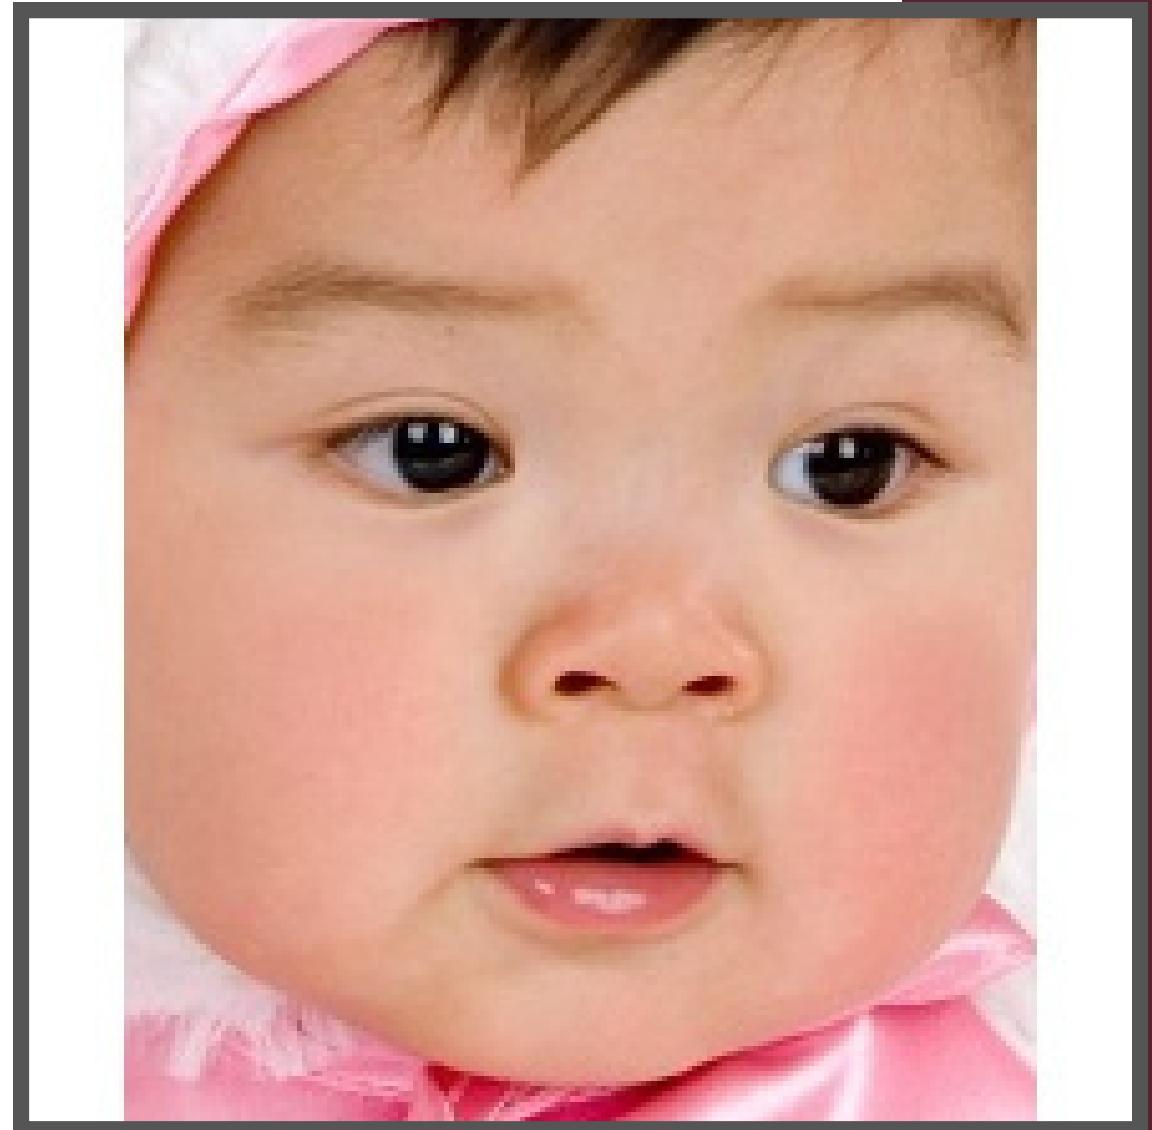
Face detection is done using library MTCNN.

# Resizing the Face Images

Resizing is needed to be able to **train** fixed-size-input CNNs.

Before resizing the images, the quality of the images were checked with certain threshold to be sure too **low-quality images are not used**.

224x224 Image



All images were resized to **224x224** resolution. White paddings are added for the empty pixels to make all the images square

# Saving in NumPy Files

All images are read and save into a numpy file for more convenient use.

The observations are saved in batches, for example 5000-observations per file.

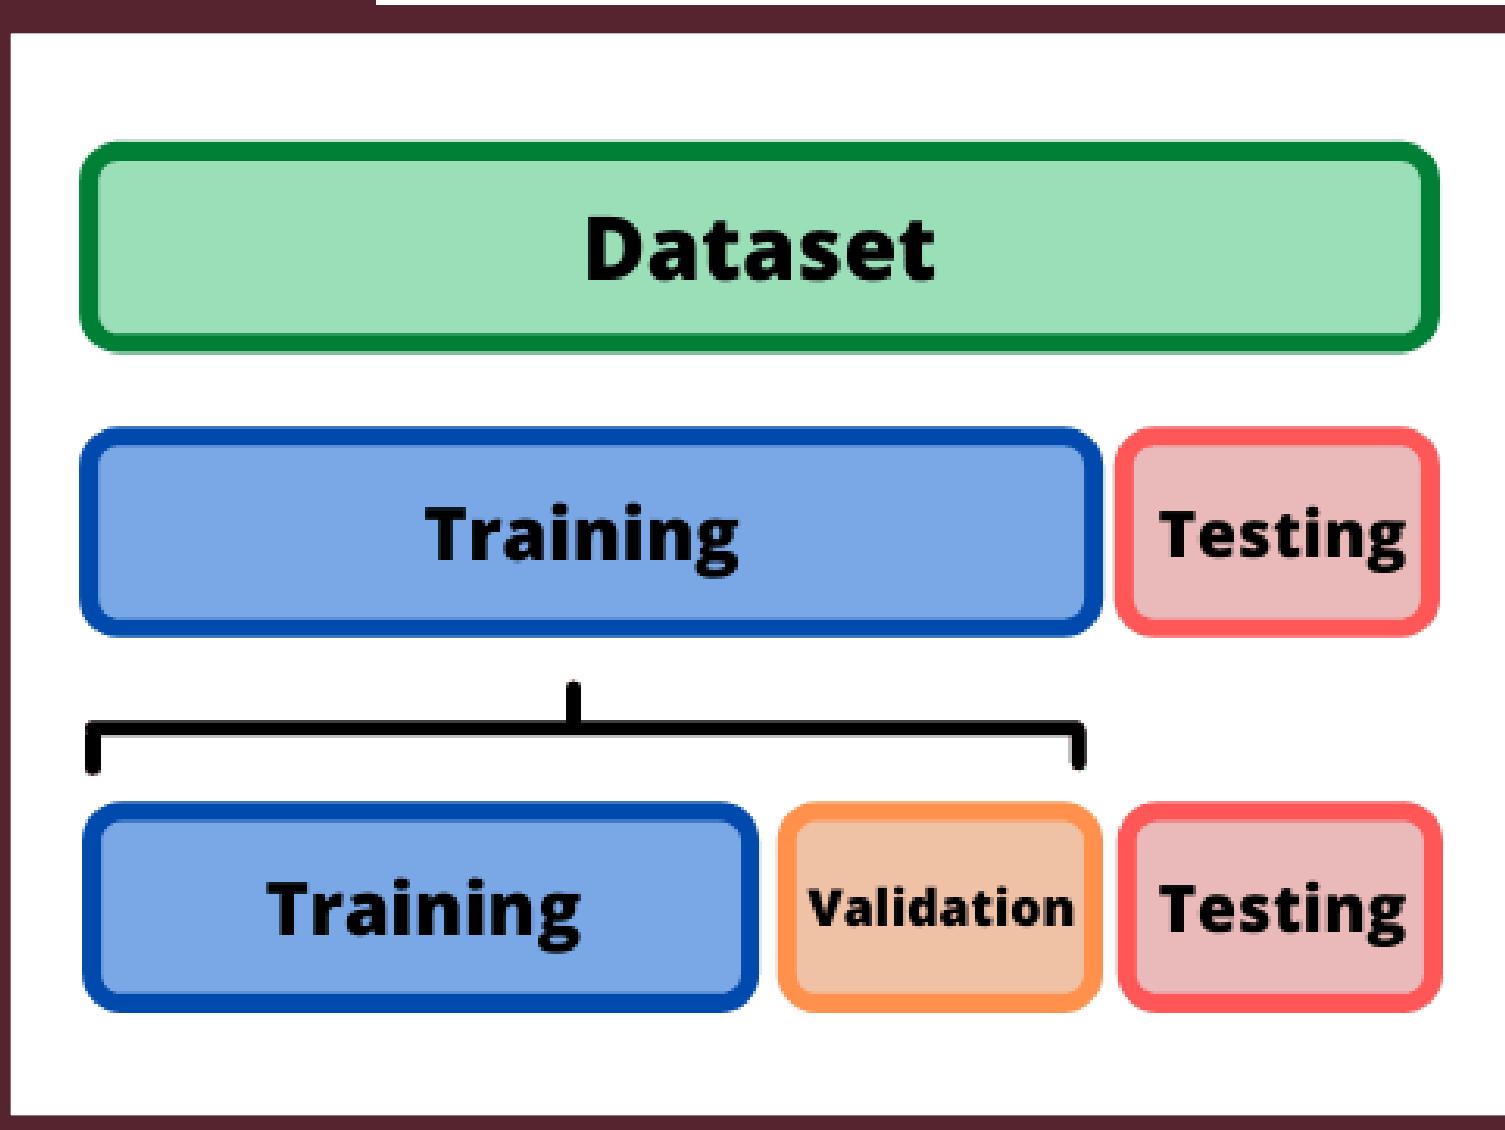
The labels are extracted and saved in a csv file.

The final total image data is of shape ()



# Train-Val-Test Split

The dataset is divided into training, validation and test sets. The data is saved in separate folders.



**Train Data** – is used for estimating the parameters of the model.

**Validation Data** – is used for model selection, hyperparameter tuning, assessing the bias-variance tradeoff.

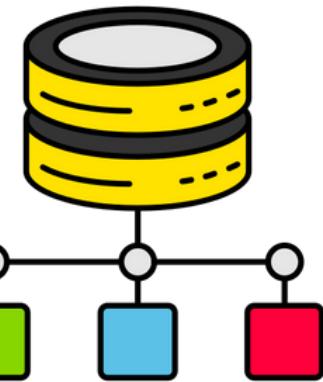
**Test Data** – is used for final evaluation of the final model.

# Problem Requirements

## Age Estimation

## Gender Prediction

**1) Creating the dataset**



# Creating the Dataset

Torch DataLoader → GetItem(Index)

**Image** (Index) with swapped dimensions

**Age** (Index)

**Gender** (Index)

```
class TorchDatasetCreator(Dataset):
    def __init__(self, dataset):
        self.dataset = dataset

    def __len__(self):
        return len(self.dataset[0])

    def __getitem__(self, idx):
        image = self.dataset[0][idx]
        image = transforms.ToTensor()(image)
        age = self.dataset[1][idx]
        gender = self.dataset[2][idx]

    return image, age, gender
```

Train Loader

Test Loader

Validation Loader

# Problem Requirements

## Age Estimation

## Gender Prediction

- 1) Creating the dataset

**2) Cost - MSE or MAE**

**2) Cost - Binary Cross-Entropy**



# Cost Function

Cost/Loss function is different for both problems, because one case is a regression problem, the other one is a classification problem.

$$L_{BCE} = -\frac{1}{n} \sum_{i=1}^n (Y_i \cdot \log \hat{Y}_i + (1 - Y_i) \cdot \log (1 - \hat{Y}_i))$$

**Classification**

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

**Regression**

# Problem Requirements

## Age Estimation

- 1) Creating the dataset
- 2) Loss – MSE or MAE

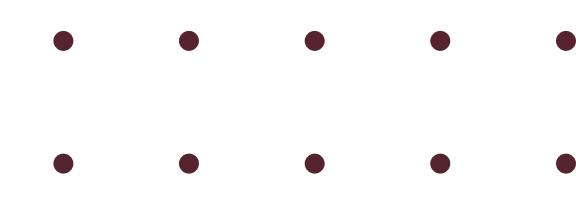
## Gender Prediction

- 2) Loss – Binary Cross-Entropy

**3) Optimizer**



# Optimizer



## Adam Optimizer

Is a go-to optimizer in many deep learning frameworks and applications due to its robust performance across a wide range of problems.

It's an extension of the basic gradient descent algorithm that incorporates several additional features to improve the efficiency and performance of the optimization process.

Easy to implement, is computationally efficient, has little memory requirements, needs less hyperparameter tuning, etc.

# Problem Requirements

## Age Estimation

1) Creating the dataset

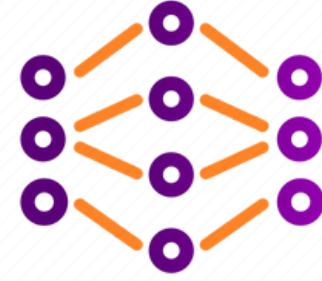
2) Loss – MSE or MAE

## Gender Prediction

2) Loss – Binary Cross-Entropy

3) Optimizer

**4) Trainer function**



# Trainer Function

**For each epoch**

Running loss = 0

**For each batch**

Predict each batch

Calculate Loss

Calculate the gradients

Optimizer updates the parameters

Zero out the accumulated gradients

Add the loss to running loss

Calculate the loss for validation data

Calculate the accuracy for validation data (**For classification**)

Print the results of the epoch

# Problem Requirements

## Age Estimation

1) Creating the dataset

2) Loss - MSE or MAE

## Gender Prediction

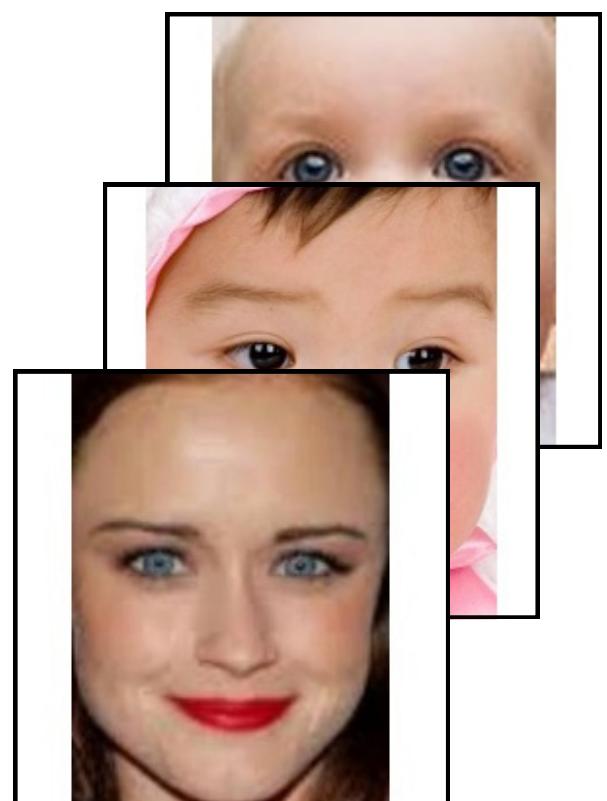
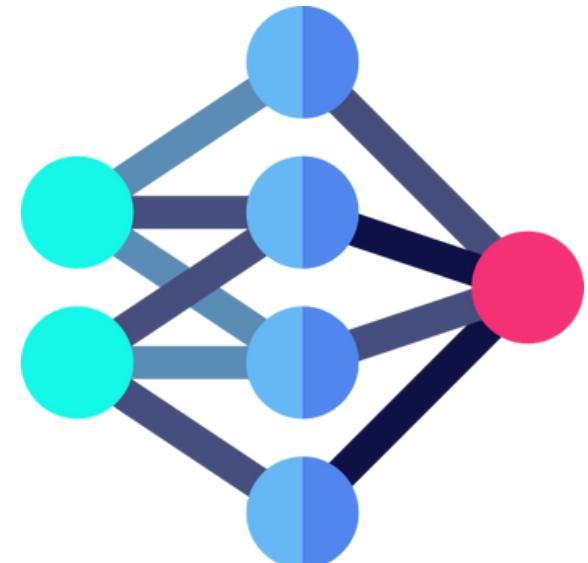
2) Loss - Binary Cross-Entropy

3) Optimizer

4) Trainer function

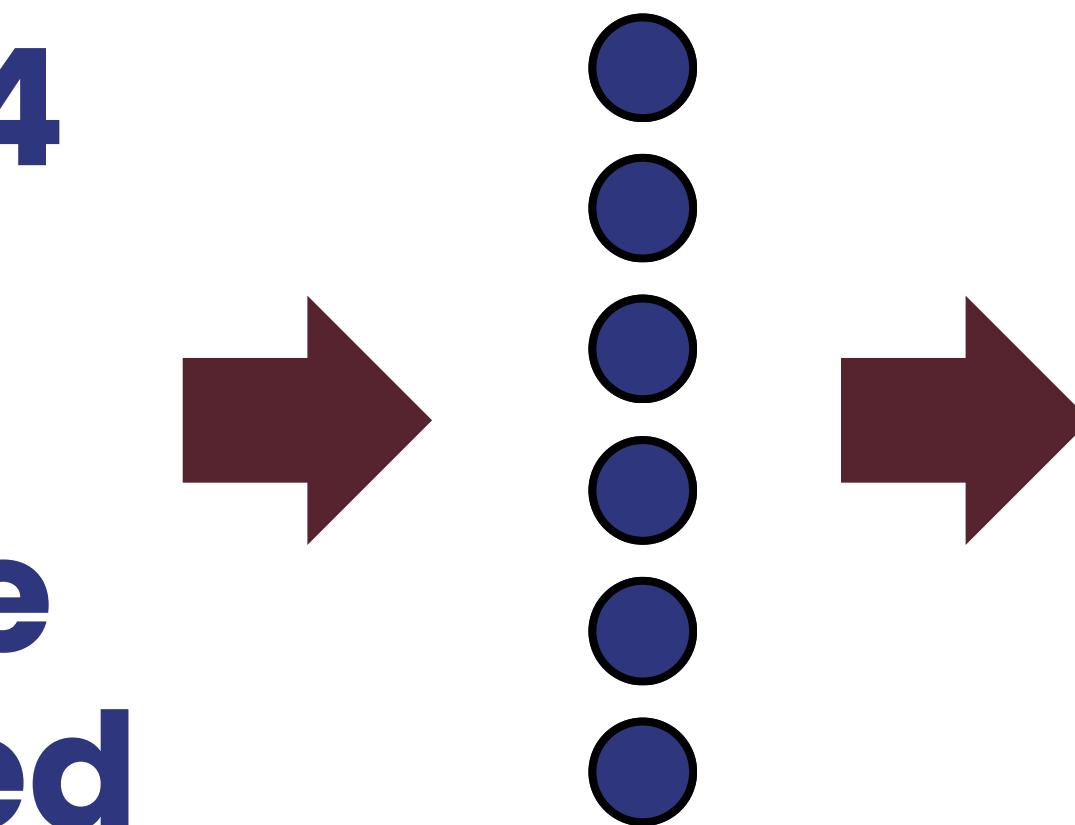
**5) CNN**

**5) CNN**



# Convolutional Neural Network

**ResNet34**  
+  
**FairFace**  
**Pretained**



Adding new FC layer at  
the end

**Sigmoid**  
**(Classification)**  
+  
**Predictions**

# Problem Requirements

## Age Estimation

1) Creating the dataset

2) Loss – MSE or MAE

## Gender Prediction

2) Loss – Binary Cross-Entropy

3) Optimizer

4) Trainer function

5) CNN

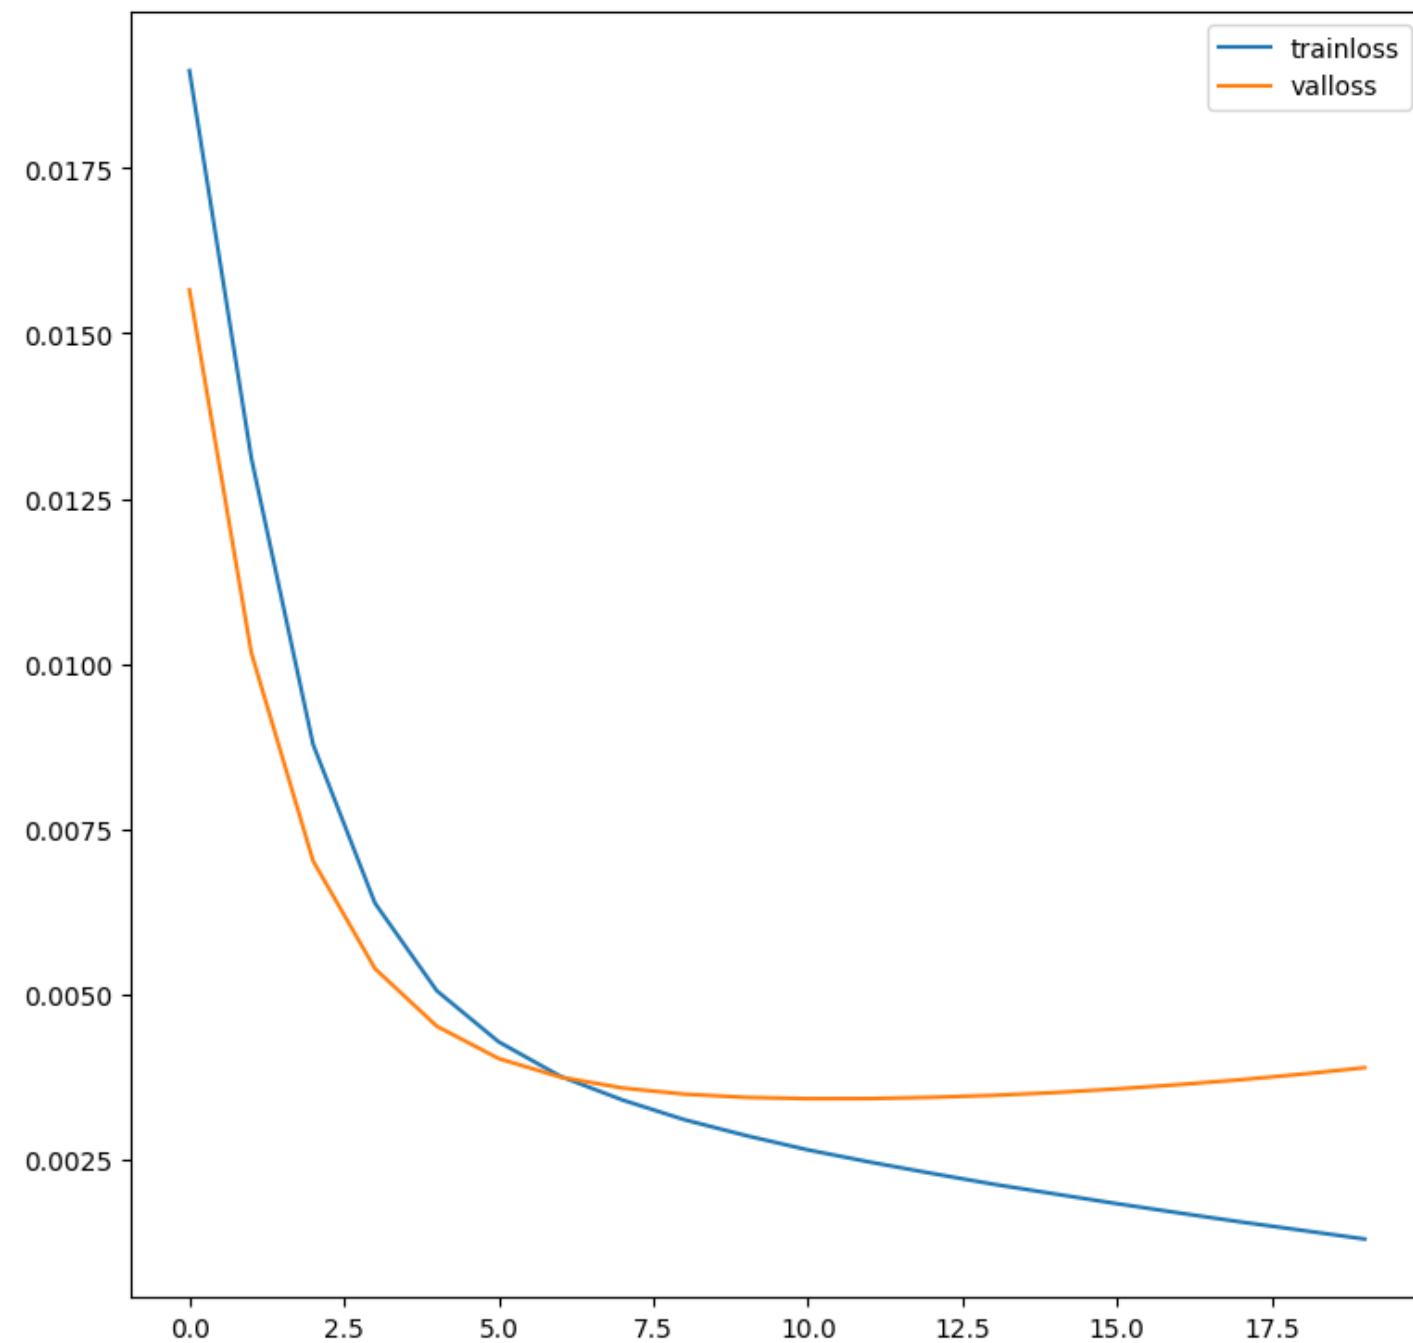
**6) Results**



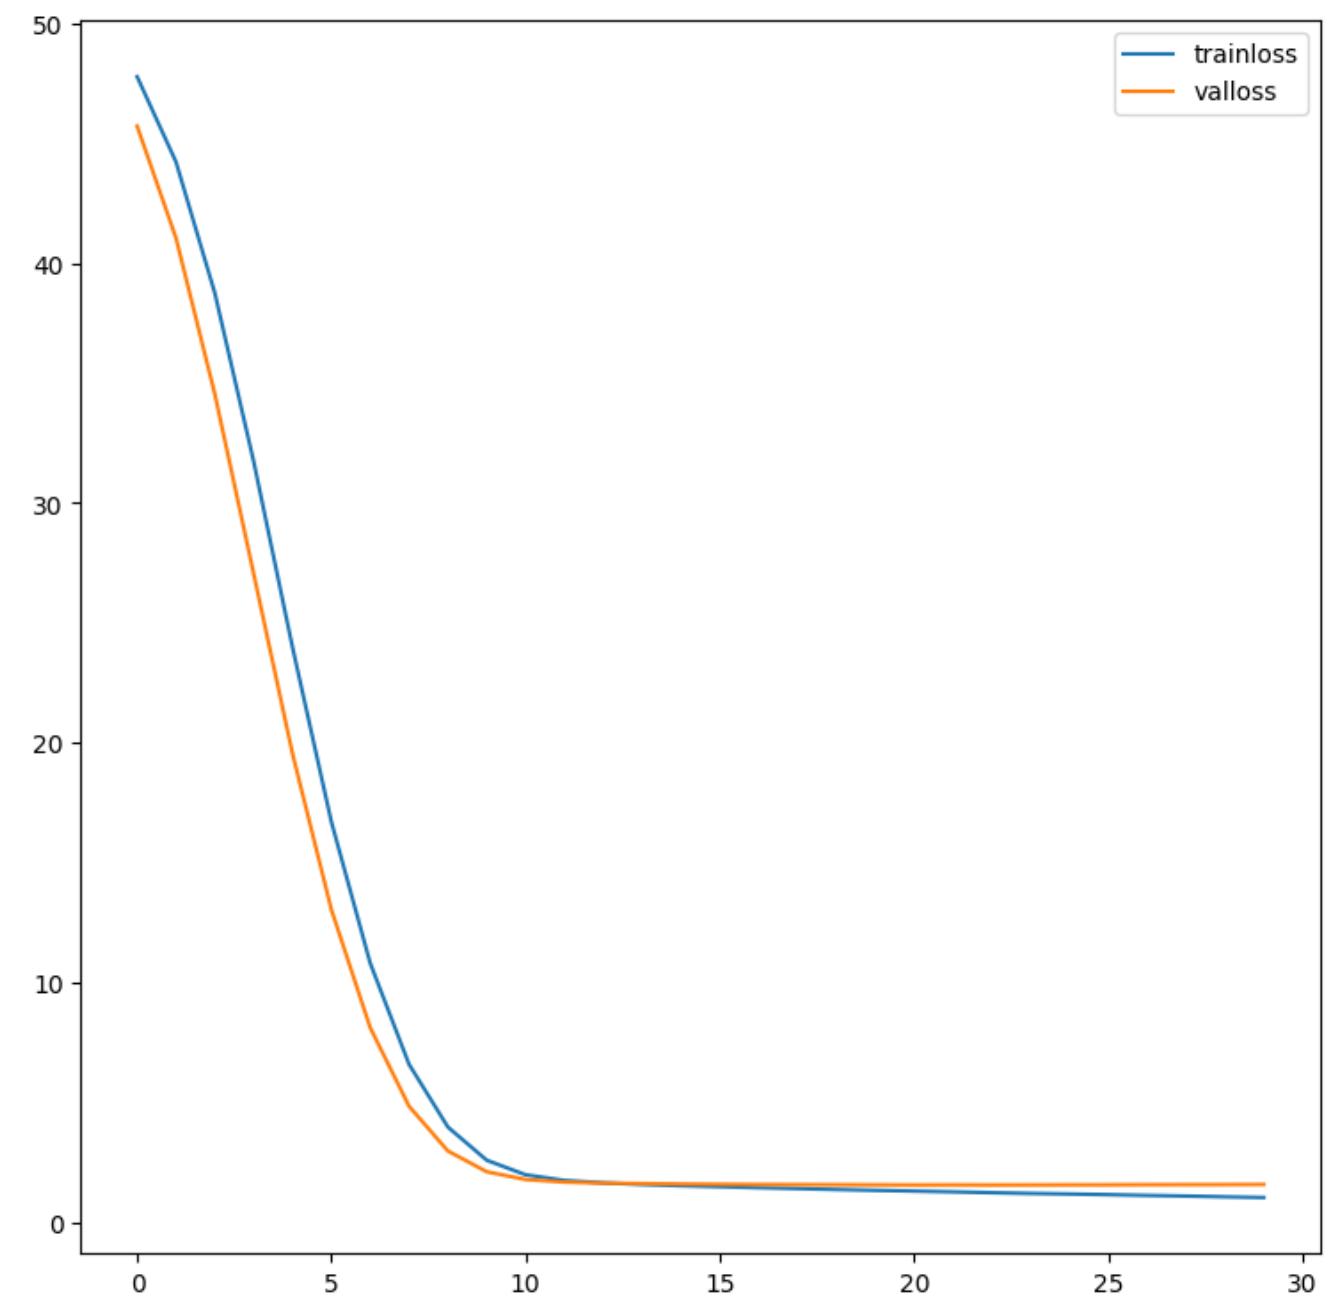
# Results

## (Learning Curves)

### Gender Prediction



### Age Prediction





# Results

## (Learning Curves)

### Gender Prediction

Train Accuracy = 97.74%

Validation Accuracy = 96.1%

Test Accuracy = 96.05%

### Age Prediction

Train Accuracy = 70.8%

Validation Accuracy = 63.4%

Test Accuracy = 63.2%

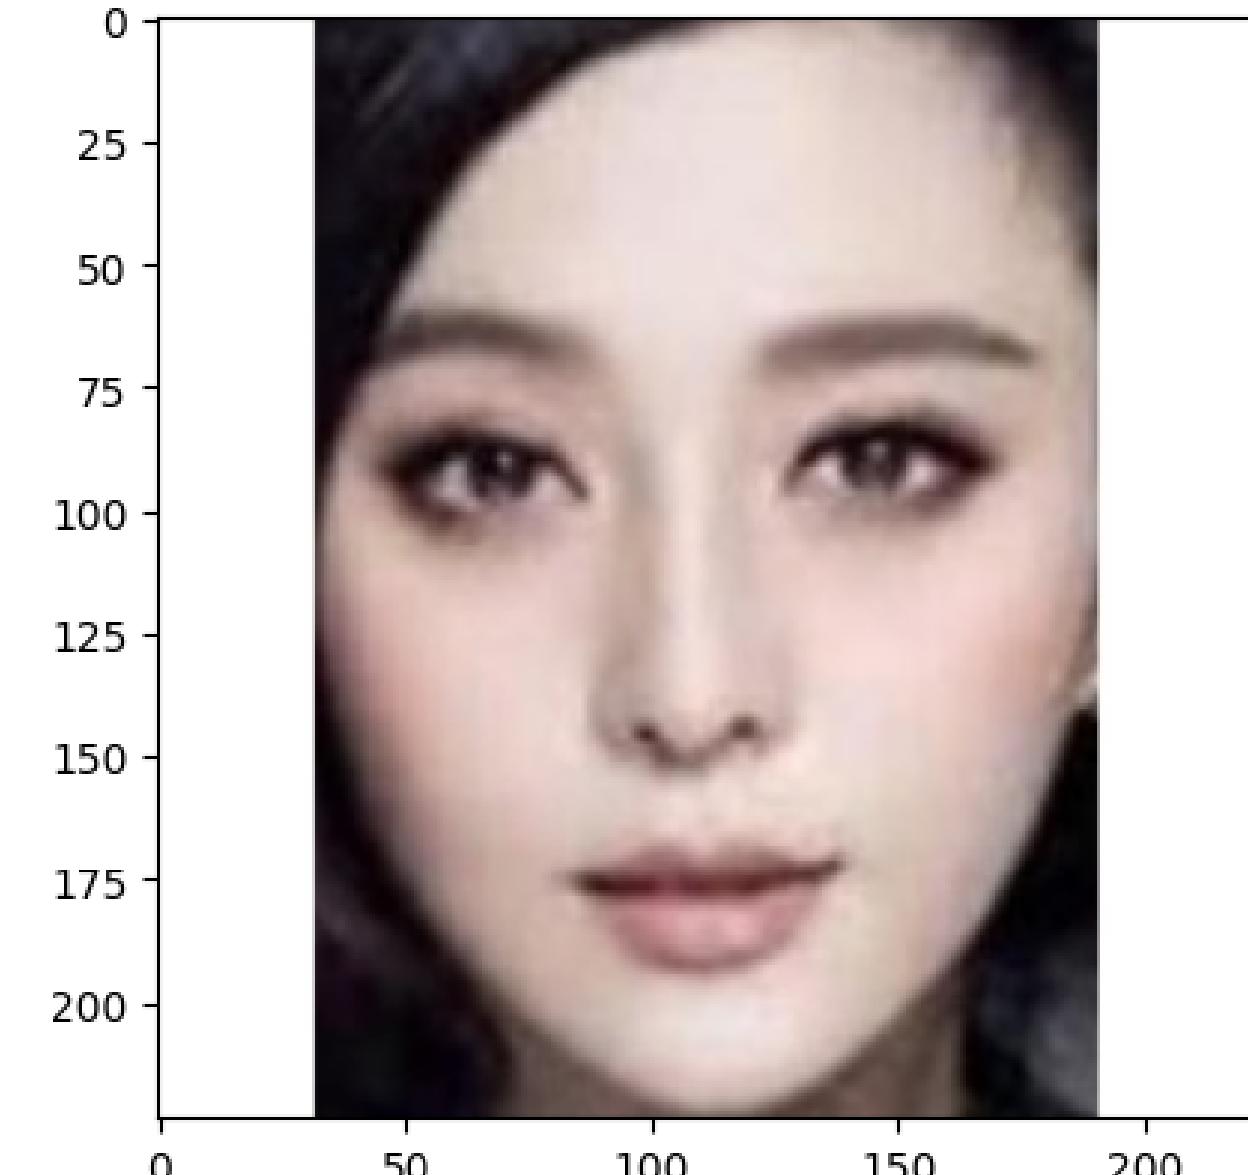
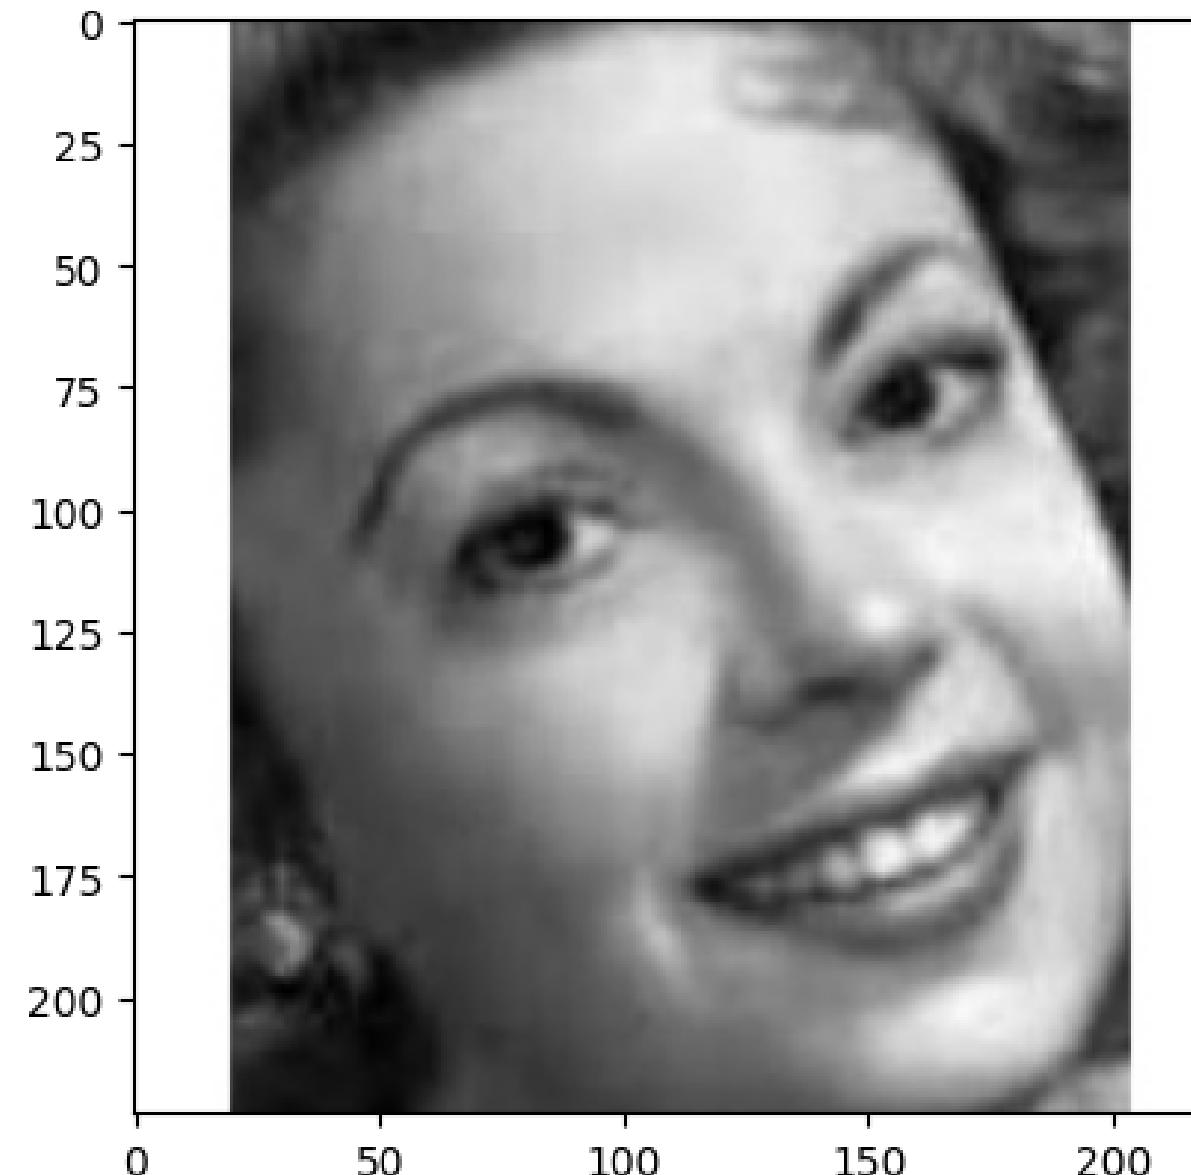
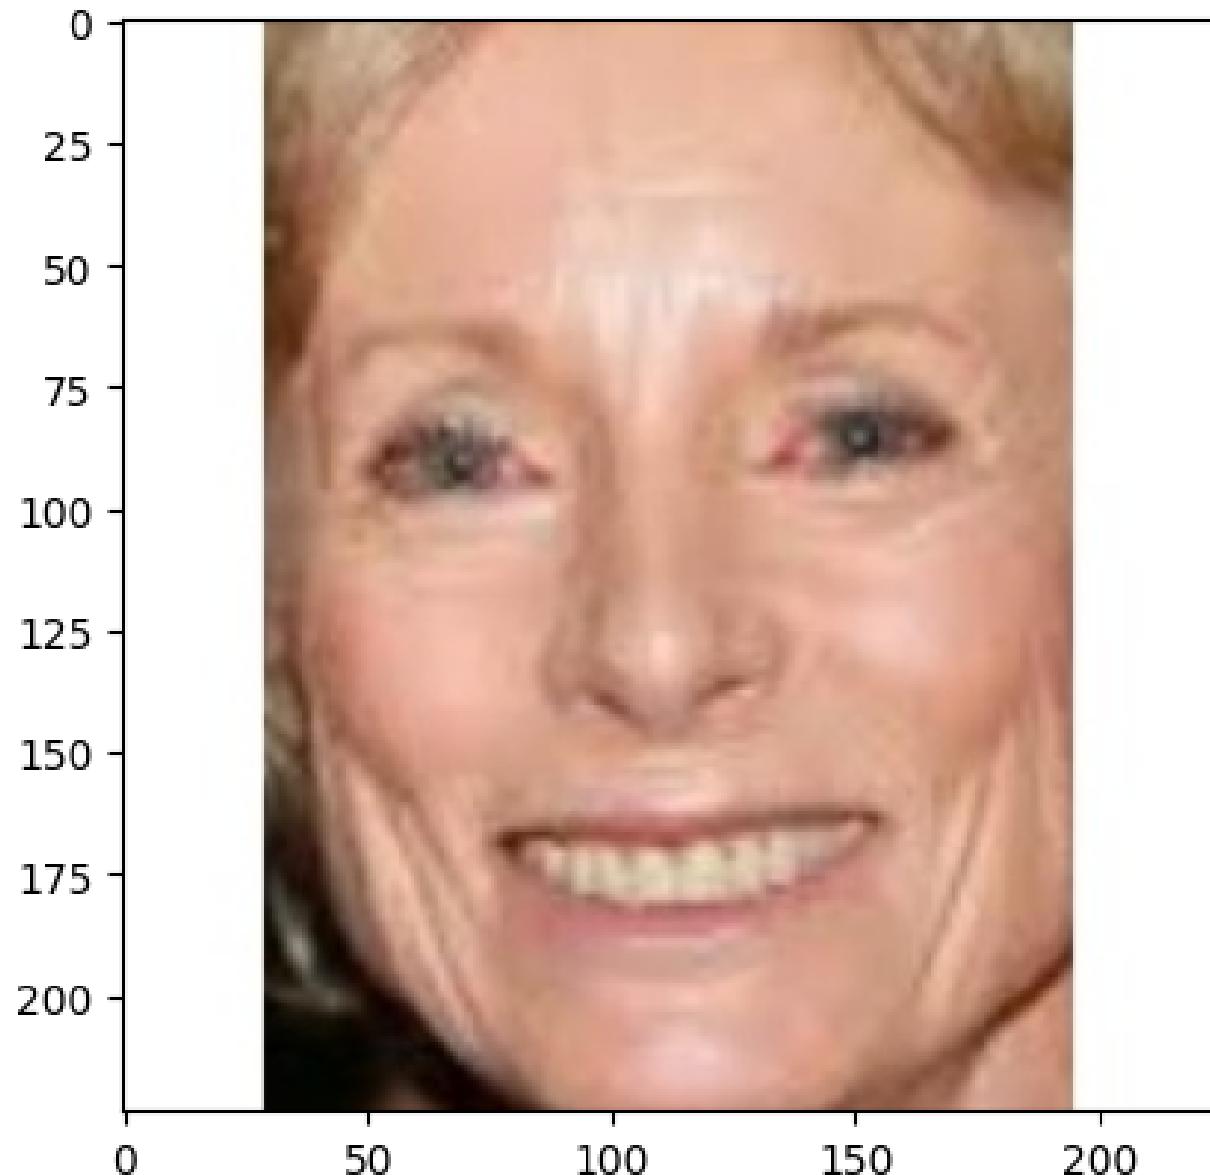
Train MAE = 4.22 years

Validation MAE = 5.24 years

Test MAE = 5.18 years



# Wrong predictions for Gender



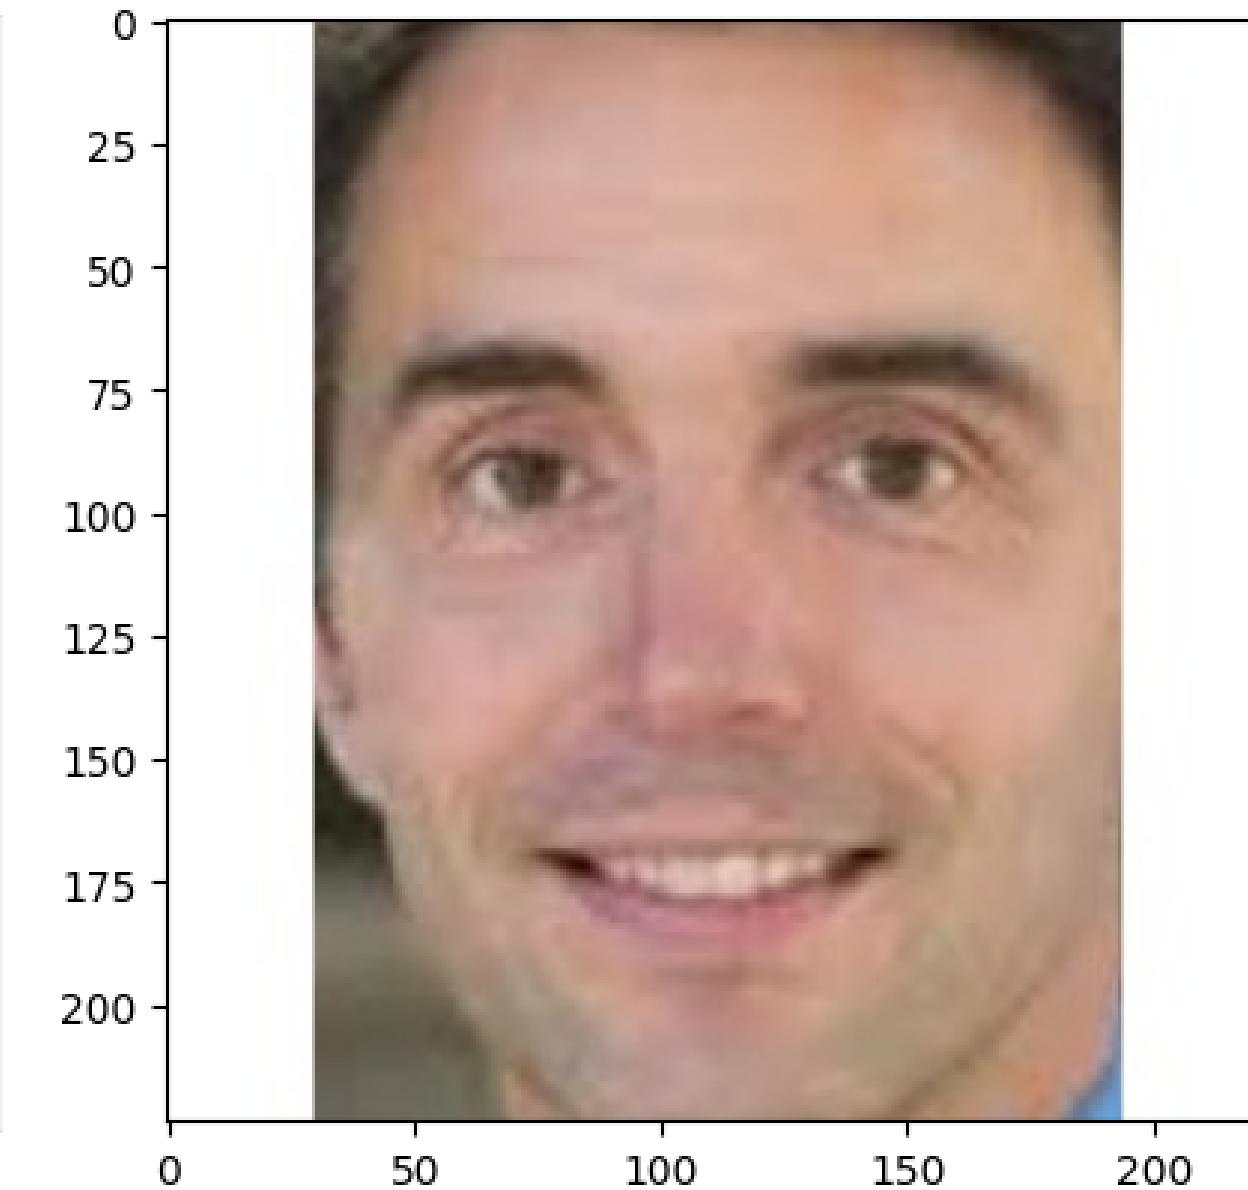
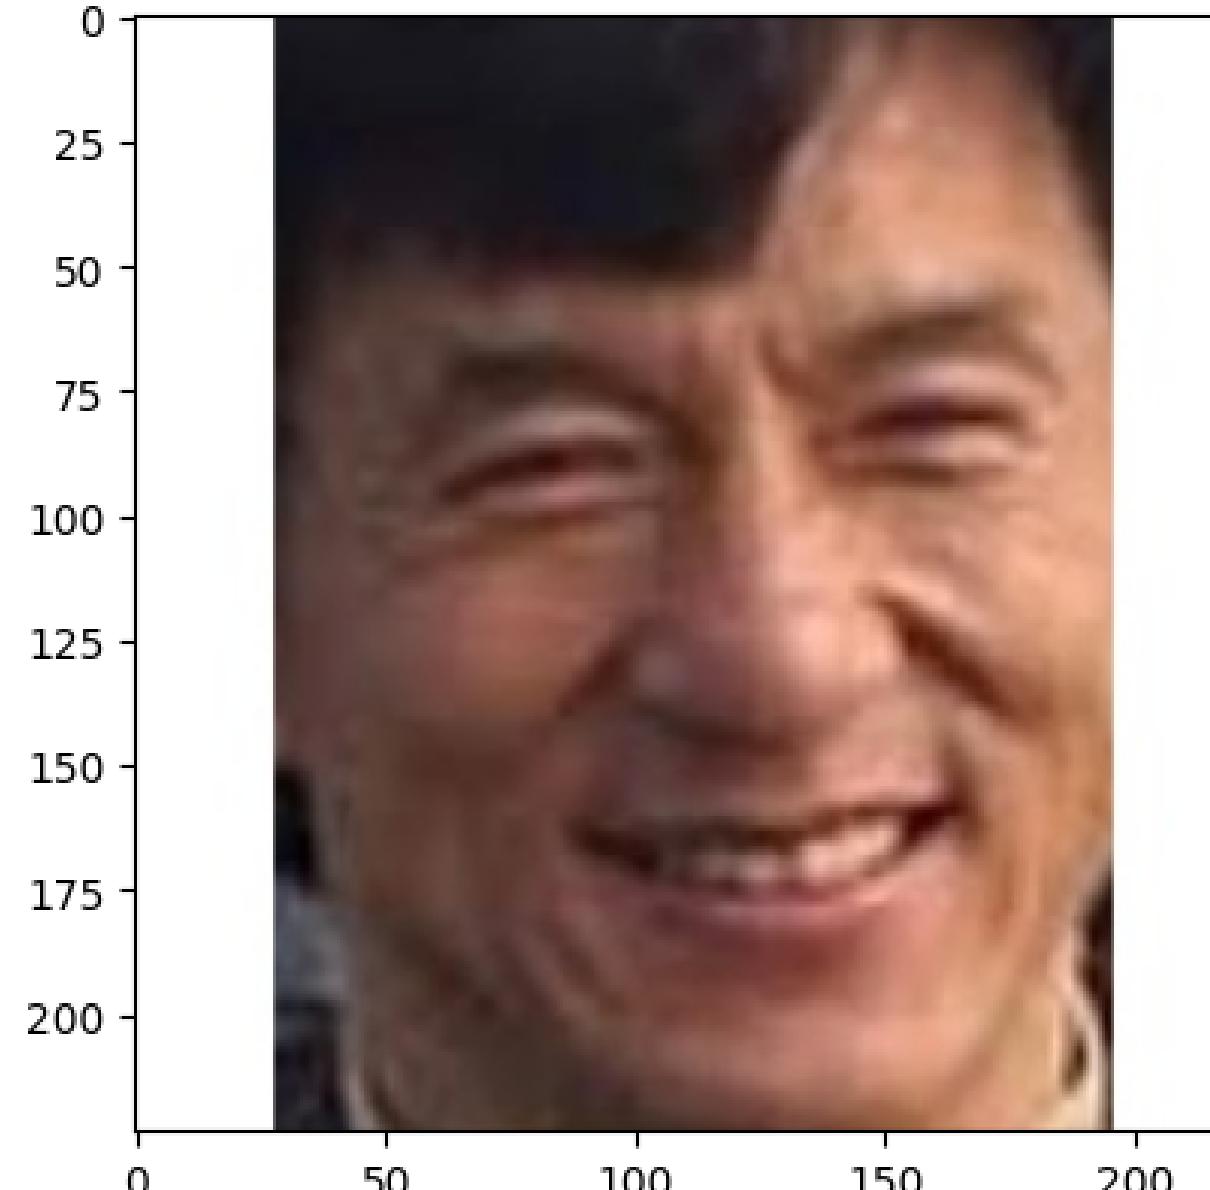
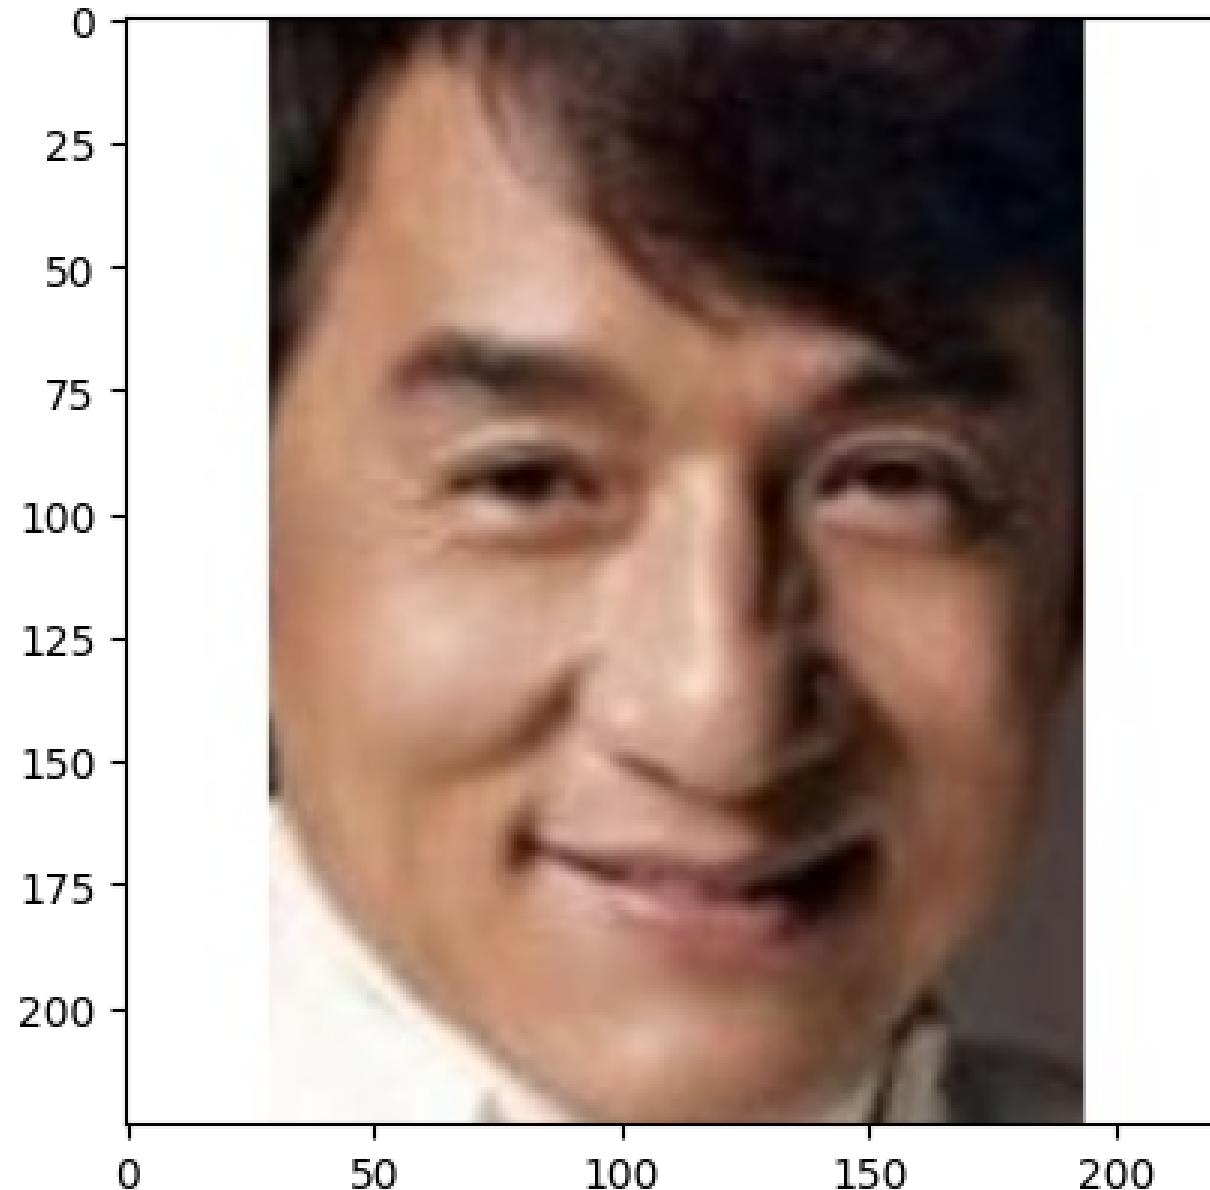
Correct label = Male  
Predicted label = Female  
Error = 0.9980271458625793

Correct label = Male  
Predicted label = Female  
Error = 0.9950774312019348

Correct label = Male  
Predicted label = Female  
Error = 0.991824209690094



# Wrong predictions for Age



---

Correct label = 59.0  
Predicted label = 35.63019561767578  
Error = 23.36980438232422

---

Correct label = 35.0  
Predicted label = 58.19916915893555  
Error = 23.199169158935547

---

Correct label = 51.0  
Predicted label = 33.14531326293945  
Error = 17.854686737060547

**Thank you!**